

CREW-AUTHORED PROJECT SUMMARY

Story Agent

An 11-agent crew that keeps the expensive model **thin** — and makes the cost of every decision visible.

Authored in parallel by the full crew · ~61k LOC TS · ~150 MCP tools · 41 skill theories · AWS Fargate · self-hosted alternative to Claude Code / Copilot.

AGENDA

What we'll cover

- 1 · Executive summary & vision — **Picard**
- 2 · Architecture — **Data**
- 3 · Cost model & ROI — **Quark**
- 4 · Security & governance — **Worf**
- 5 · Infrastructure & deployment — **Geordi**
- 6 · DevOps & reliability — **O'Brien**
- 7 · Quality & go-live — **Yar**
- 8 · Client value — **Troi**
- 9 · Roadmap — **Riker**
- 10 · System health & honesty — **Crusher**
- 11 · Narrative & close — **Uhura**

Redefining the coding assistant

- A self-hosted, 11-agent orchestration system for cost-effective code generation on OpenRouter's multi-provider pool.
- The premium model only orchestrates — the cheap crew does the token-heavy work.
- Built on the Model Context Protocol (open standard) — the same crew plugs into Claude Code, Cursor, Copilot.
- WorfGate governance + pgvector RAG give it security and compounding memory.

Client value: Scalable, secure, cost-attributed AI coding — faster, safer development with measurable ROI.

Control flow: how the system thinks



- A request enters over MCP — the open standard, so any AI client can drive the same crew.
- The orchestrator hands off to the mission pipeline (deliberation) or the agent-core loop (coding).
- Quark routing dispatches each task to the cheapest adequate model in a multi-provider pool.
- WorfGate wraps every operation; RAG recalls before acting and stores after.

Client value: A self-improving, cost-attributed assistant with automatic security guardrails that connects to the client's existing tools.

Cost model & ROI

- Thin premium orchestration — Anthropic is used only for tier-4 (architecture / security).
- Most work runs on DeepSeek / Llama; a full crew deliberation costs ~\$0.002–\$0.003.
- Quark routing picks the cheapest model that meets each task's required capability tier.
- The control-lane ledger measures crew-vs-premium spend, so savings are reported, not promised.

Client value: Significant, measurable savings vs running a frontier model for everything — with the split shown on a ledger.



SECTION 4 **WORF · SECURITY**

Security & governance: WorfGate

- Credential broker: secrets are injected at runtime and never written to logs, prompts, or memory.
- Green / yellow / red local governor: risky filesystem & shell ops are auto-remediated (paths clamped, --force downgraded), not just blocked.
- Audit trail by design: every crew deliberation is cost-attributed and every autonomous action is recorded.
- Self-hosted — controlled data never leaves the client's environment.

Client value: Regulated clients get bounded autonomy: credentials never exposed, every action audited and policy-gated.

SECTION 5 GEORDI · INFRASTRUCTURE

Infrastructure & deployment

- AWS Fargate deployment via Terraform + Docker — scalable, reproducible infrastructure.
- Multi-provider model reachability (Meta / OpenAI / DeepSeek / Anthropic) through OpenRouter.
- WorfGate's local governor mitigates risky operations at runtime.

Client value: Cost-efficient, reliable, secure infrastructure that runs in the real environment, not just a demo.

DevOps & reliability

- Operability first: runbooks + WorfGate auto-remediation keep the system runnable by someone who wasn't in the room.
- Live in-flight progress via the new async-status registry — surfaced on every prompt.
- Timeout-as-terminal-state guardrail: a silent hang now shows up instead of wasting time invisibly.
- Verified deploy on AWS Fargate (Terraform + Docker).

Client value: Less downtime and no silent stalls — clients see exactly what the system is doing as it does it.

SECTION 7 YAR · QUALITY

Quality & go-live criteria


- Evidence-based signoff — no “seems fine” approvals; changes ship with tests.
- A documented shadow test is the honest go/no-go for full autonomy: it is a criterion, not yet passed.
- WorfGate clamps risky ops before they escalate; async-status + timeouts bound live risk.
- Cost-aware quality: deliberation micro-costs stay predictable under multi-provider routing.

Client value: Clients get compounding improvement without unbounded cost risk or unvetted autonomy.

Client value & experience

- Cost you can see — transparent, per-decision attribution.
- Governed, secure operations by default (WorfGate + Quark routing).
- A compounding knowledge base (RAG) that carries across sessions and across AI tools.
- No vendor lock-in — the crew brain is portable via MCP.

Client value: Faster, cheaper development with governance and memory the client owns — not the vendor.



SECTION 9 RIKER · IMPLEMENTATION

Roadmap — what's next

- Crew-aware feature branches tied to Aha stories — the whole crew knows the active branch/story.
- Async-visualization follow-up: streaming (Redis→SSE) + web and VS Code renderers, full producer wiring.
- Fix the crew-table write path (RLS/permission) so DB-backed personas/skills persist.
- Run the shadow test and record the result — turn the go-criterion into evidence.

Client value: Clearer progress visibility, tighter PM integration, and a documented path to full autonomy.

System health & honest current state

- Working today: crew deliberation, the coding loop, cost routing, WorfGate, RAG, async-status, Fargate deploy.
- Known gaps: shadow test not yet passed; a crew-table write path needs fixing; cost figures are per-deliberation, not audited aggregates.
- No verified production customers yet — reference clients are illustrative.
- Monitoring baseline + candor are treated as first-class, not afterthoughts.

Client value: Clients can trust the claims because the limits are stated as plainly as the strengths.

How to tell this story

- Lead with the economics: the expensive model does almost none of the work.
- Show the control flow once — then let cost, governance, and memory hang off it.
- Close on candor: it runs today; the bar for full autonomy is defined and not yet met.
- Every claim here is grounded in the codebase, with caveats kept visible.

Client value: A clear, credible narrative a decision-maker can repeat to someone who wasn't in the room.

CLOSE

“Keep the expensive model thin; make the cost of every decision visible.”

It runs today. The shadow test that gates full autonomy is defined — and not yet cleared.

Self-hosted · governed · cost-attributed · portable across AI tools. Every figure verified against the codebase; caveats kept honest.